

Govern every AI interaction before it acts.


A provider-independent control layer for model,
agent, tool and API traffic.

Built locally. Enforced centrally. Explained

ALLOW

REDACT

BLOCK



AI traffic is becoming a new production boundary.

Every prompt, tool call and model response can carry sensitive data or risky instructions. The control point must sit between intent and execution.

Sensitive data

PII, credentials and private-key material can move into model or tool paths.

Unsafe intent

Prompt injection, instruction override and exfiltration intent need policy decisions.

Blind operations

Without safe evidence, operators cannot explain what happened or scale trust.

The opportunity

Make security a reusable service between every caller and every target.

One control point. One decision model.

The system normalizes an interaction, evaluates registered controls, applies central policy, records safe evidence and forwards only what is eligible.



Every decision is structured, explainable and auditable before eligible dispatch.

A focused security pack that already works end to end.

The current build is deliberately bounded: it proves the control-layer pattern with local targets, strict startup policy and deterministic tests.

6

registered controls

email, bearer, PEM, GitHub, labelled SSN, known attacks

3

central actions

ALLOW preserves, REDACT transforms, BLOCK stops

2

target boundaries

local echo and operator-configured loopback Ollama

19

demo scenarios

benign, deterministic, semantic and failure modes

The system is provider-independent by construction:

controls return findings; policy owns enforcement; adapters own protocol behavior; targets remain replaceable.

A simple outcome users can trust.

The decision engine makes the security posture visible in one line, while retaining the structured reasons behind it.



ALLOW

No finding maps to a stronger action.

Forward the original interaction exactly once.



REDACT

A supported PII finding needs protection.

Replace selected spans with [REDACTED] before forwarding.



BLOCK

A credential or known attack indicator is detected.

Make zero target calls and return a safe reason.

Central precedence: **BLOCK > REDACT > ALLOW**

Protection with honest boundaries.

The pack recognizes precise supported shapes and keeps its claims narrow. That makes the behavior testable, explainable and safe to extend.

CONTROL	FINDING FAMILY	DEFAULT	BOUNDARY
Email address	pii.email	REDACT	bounded ASCII dot-atom
Labelled US SSN	pii.us_ssn	REDACT	labelled, fictional syntax
Bearer / PEM / GitHub	secret.*	BLOCK	supported token/key shapes
Known attack literals	attack.*	BLOCK	5 exact inert signatures
Semantic security	semantic.*	BLOCK	optional local evaluator

Current limits: no output inspection, auth, budgets, policy reload, remote provider or general secret / jailbreak detection. Those are product expansion points, not hidden claims.

Evidence that helps operators move fast.

Every relevant decision becomes safe telemetry. The dashboard turns that evidence into a living view of posture, timing and outcomes.

AUDIT EVENT

server interaction ID

UTC timestamp

target + policy digest

evaluated controls + enablement

action + finding codes

evaluation / invocation / total timing

DASHBOARD

ALLOW

posture

REDACT

protected

BLOCK

stopped

UNKNOWN


evidence gap

Privacy by design

Audit and reporting retain safe metadata, never raw prompt content, matched values, spans, hashes, scores or exception text.

From live walkthrough to repeatable evidence.

The local dashboard combines a read-only security overview with an interaction workbench for custom input and prepared scenarios.




19 code-owned scenarios with expected explanations and actual-result reporting

The operator can see what was expected, what happened, why policy chose the action and whether the target completed. A mismatch stays visible instead of being rewritten as success.

Scale the boundary first; scale the adapters next.

The architecture keeps the security decision core small and composable. That gives the product a credible path from local proof to platform service.



What scales cleanly
stateless controls | immutable registrations | policy-owned decisions | replaceable target adapters

What remains intentionally local today: synchronous request handling, loopback Ollama, single-node SQLite and no distributed queue. These are clear deployment boundaries, not architecture dead ends.

WHY THIS WINS

A control layer people can adopt before they trust it at scale.

It starts with a narrow, testable promise: no eligible AI interaction reaches a target without a policy decision and safe evidence.

1,614

unit + integration checks

9

Chromium journeys

PASS

fresh correctness + security reviews

0

target calls after BLOCK

Protect the interaction. Preserve the evidence. Earn the next integration.

AI CONTROL LAYER / local proof today / scalable policy boundary tomorrow