

Govern every AI interaction before it acts.


A provider-independent control layer for model, agent, tool and API traffic.

Built locally. • Enforced centrally. • Explained clearly.



EVERY INTERACTION

ONE DECISION



AI traffic is the new production boundary.

Every prompt, tool call and model response can carry sensitive data or risky instructions. The control point must sit between intent and execution.



01

Sensitive data

PII, credentials and private-key material can move into model or tool paths.



02

Unsafe intent

Prompt injection, instruction override and exfiltration intent need policy decisions.



03

Blind operations

Without safe evidence, operators cannot explain what happened or scale trust.


THE OPPORTUNITY



**Make security a reusable service
between every caller and every target.**

One control point. One decision model.

The system normalizes an interaction, evaluates registered controls, applies central policy, records safe evidence and forwards only what is eligible.



Every decision is structured, explainable and auditable before eligible dispatch.

A focused security pack, working end to end.

The current build is deliberately bounded: it proves the control-layer pattern with local targets, strict startup policy and deterministic tests.

6



registered controls

email, bearer, PEM, GitHub,
labelled SSN, known attacks

3



central actions

ALLOW preserves, REDACT
transforms, BLOCK stops

2



target boundaries

local echo and operator-
configured loopback Ollama

19



demo scenarios

benign, deterministic, semantic
and failure modes

Provider-independent by construction



Controls return findings



Policy owns enforcement



Adapters own protocol behavior



Targets remain replaceable

A simple outcome users can trust.

The decision engine makes the security posture visible in one line, while retaining the structured reasons behind it.



ALLOW

WHEN

No finding maps to a stronger action.

THEN

Forward the original interaction exactly once.



REDACT

WHEN

A supported PII finding needs protection.

THEN

Replace selected spans with [REDACTED] before forwarding.



BLOCK

WHEN

A credential or known attack indicator is detected.

THEN

Make zero target calls and return a safe reason.



Central precedence

BLOCK



REDACT



ALLOW

the stronger action always wins



Protection with honest boundaries.

The pack recognizes precise supported shapes and keeps its claims narrow. That makes the behavior testable, explainable and safe to extend.

CONTROL	FINDING FAMILY	DEFAULT	BOUNDARY
Email address	pii.email	REDACT	bounded ASCII dot-atom
Labelled US SSN	pii.us_ssn	REDACT	labelled, fictional syntax
Bearer / PEM / GitHub	secret.*	BLOCK	supported token / key shapes
Known attack literals	attack.*	BLOCK	5 exact inert signatures
Semantic security	semantic.*	BLOCK	optional local evaluator



Current limits

Deliberately not claimed by today's build:

output inspection

auth

budgets

policy reload

remote provider

general secret / jailbreak detection



Product expansion points, not hidden claims.



Evidence that helps operators move fast.


Every relevant decision becomes safe telemetry.
The dashboard turns that evidence into a living view of posture, timing and outcomes.



Audit event

structured record per decision

- ✓ server interaction ID
- ✓ UTC timestamp
- ✓ target + policy digest
- ✓ evaluated controls + enablement
- ✓ action + finding codes
- ✓ evaluation / invocation / total timing



Privacy by design

Audit and reporting retain safe metadata only. They never keep:

× raw prompt content

× matched values

× spans

× hashes

× scores

× exception text

From live walkthrough to repeatable evidence.

The local dashboard combines a read-only security overview with an interaction workbench for custom input and prepared scenarios.



BENIGN

ALLOW

safe request



DETERMINISTIC

REDACT / BLOCK

PII, token, key, attack



SEMANTIC

BLOCK

injection, override, exfiltration




FAILURE

EXPLAIN

evaluator / target unavailable

19

code-owned scenarios with
expected explanations and
actual-result reporting




The operator can see what was expected, what happened, why policy chose the action and whether the target completed.

**A mismatch stays visible instead
of being rewritten as success.**

Scale the boundary first; scale the adapters next.

The architecture keeps the security decision core small and composable. That gives the product a credible path from local proof to platform service.



Intentionally local today: synchronous request handling, loopback Ollama, single-node SQLite and no distributed queue. These are clear deployment boundaries, not architecture dead ends.

A control layer people can adopt before they trust it at scale.

Thank
you!

It starts with a narrow, testable promise: no eligible AI interaction reaches a target without a policy decision and safe evidence.

1,614



unit + integration checks

9



Chromium journeys

PASS



fresh correctness +
security reviews

0



target calls after BLOCK



Protect the interaction.



Preserve the evidence.



Earn the next integration.

LOCAL PROOF TODAY · SCALABLE POLICY BOUNDARY TOMORROW

